



UNIVERSITY OF SOUTHERN DENMARK
DEPARTMENT OF MATHEMATICS AND COMPUTER SCIENCE,
IMADA

MASTER THESIS

Building a Package Manager for Jolie

Author:

Dan Sebastian THRANE

Supervisor:

Fabrizio MONTESI

May 24, 2017

Contents

1	Introduction	4
2	Background	5
2.1	Microservices	5
2.2	Package Managers	5
2.3	Introduction to Jolie	5
2.4	The Jolie Engine and Interpreter	10
2.5	A Complete Application with Jolie	12
3	Jolie Modules	18
3.1	Summarizing the Problem	18
3.2	Modules: Design and Implementation	18
3.3	Configuration	19
3.4	Implementing the Configuration Format	30
3.5	Examples	39
4	Jolie Packages	45
4.1	Introduction	45
4.2	Architecture	46
4.3	Package Manifest	46
4.4	Dependencies	50
4.5	Lock Files	50
4.6	Lifetime Hooks	52
4.7	The <code>.pkg</code> Format	53
4.8	Integrity Checks	54
5	Package Manager	55
5.1	Introduction	55
5.2	Registry	56
5.3	Security	57

5.4	Calculating Dependency Trees	61
5.5	The JPM Back End	61
5.6	Authentication Tokens	63
5.7	The JPM Cache	63
5.8	The Command Line Interface	64
5.9	Examples and Discussion	66
6	Conclusion	69
A	Appendix	70
A.1	JPM Manifest Specification	70
A.2	COL Grammar	74
A.3	Scripts from Lifetime Hooks Example	75

Introduction

Introduction goes here

Background

Background goes here.

2.1 Microservices

- General stuff about microservices, actually introduce the thing

2.2 Package Managers

Some background on package managers.

2.3 Introduction to Jolie

Jolie is a service-oriented programming language, and is build to support a microservice natively. In this section we will cover what kind of language Jolie is, and how it is currently used.

Jolie has a C-inspired syntax, and is dynamically typed. Its interpreter is written in Java.

The language has no native functions or methods, but instead works in processes. A process has no arguments, and does not contain any stack (in the case of recursive calls). There are two pre-defined processes, which will always be called by the interpreter, these are called `init` and `main`.

```
1  include "console.iol"
2
3  define PrintOutput {
4      println@Console(output)() // Prints 'OK'
5  }
6
7  init { a = 1 }
8
9  main {
10     b = 2;
11     c = a + b; // c = 3
12     if (c == 3) {
13         output = "OK"
14     } else {
15         output = "Bad"
16     };
17     PrintOutput // Calls the defined process 'PrintOutput'
18 }
```

Listing 1: A very simple Jolie program

Listing 1 shows a very simple programming language, in what looks like what you might expect from a dynamic language with C-inspired syntax. However a few things may also strike you as odd.

First of all there are typos on lines 12, 14 or 16, the semicolon is not needed here, in fact it would be a syntax error. The reason for this is that the semicolon isn't used strictly for parsing purposes, but it instead for having multiple statements in a process. The "semicolon" statement, also called a sequence statement, has a syntax of `A ; B`, which should be read as: first perform statement A, then perform B. The sequence statement requires both of the operands to be present, hence the syntax error. Another similar statement is the parallel statement, which has a syntax of `A | B`, which reads as: do A and B in parallel. Using these operators together allows the programmer to easily create a fork-join workflow. This is typically used in microservices when we want to collect data in parallel, and continue once all of the data has been retrieved.

Secondly has slightly different rules for scoping. In Jolie everything not defined in the global scope goes into the same scope. This also persists through calls to defines. This is the reason that `PrintOutput` can use the output variable.

Several execution modes exists. The default execution mode, which was used in Listing 1 is `single`. This means that the `main` process is run just a single time. Two more modes exists, those being `concurrent` and `sequential`. TODO Some more stuff

Ports are the primitive that Jolie uses for communication, two types of ports exists: input and output. Ports describe a running service, where it is located (**Location**), and how to speak to it (**Protocol**), and finally which operations it supports (**Interfaces**). In Listing 2 we see a simple output port which contacts `example.com` on port 42000, using `http`. Note that it is only the ports that deal with the protocol, everything inside of the code is completely agnostic with respect to the protocol being used for communication. As a result it is easy to change a service from communicating using one protocol to another.

```
1 outputPort Example {
2     Location: "socket://example.com:42000"
3     Protocol: http
4     Interfaces: IExample
5 }
```

Listing 2: A simple output port which contacts the Google website

The interface which the port uses is called `IExample`, a full definition of it can be seen in Listing 3. Two types of operations exists in Jolie, namely `RequestResponse` and `OneWay`. The difference being fairly self-explanatory, the first receives a request and returns a response, the other simply receives a request, and produces no response.

```
1 interface IExample {
2     RequestResponse:
3         anOperation(RequestType) (ResponseType)
4
5     OneWay:
6         hello>HelloType)
7 }
```

Listing 3: An interface in Jolie defines which operations a port exposes


Whenever the Jolie interpreter invokes an operation on an output port, or receives a request on an input port, the types will be checked. This check ensures that we don't send out incorrect requests, and ensures that we do not attempt to process an incorrect request. Listing 4 show the request and response type of `anOperation`.

```
1  type RequestType: void {
2      .a: int
3      .b: string
4      .c: bool
5      .d: double
6      .e: any // any primitive type
7  }
8
9  type ResponseType: int {
10     .aFixedArray[1, 3]: string
11     .aNonFixedArray[0, *]: string
12     .fieldWithChildren: void {
13         .a: void {
14             .b: int
15         }
16     }
17 }
```

Listing 4: Jolie types are tree-like structures

In Jolie types are tree-like structures, very similar to how, for example, XML would be represented. Importantly the root may also contain a value, this is different from how most other programming languages work. This also means that certain encodings may have problems with this. JSON a popular format for serialization does not support root values, as a result Jolie will encode the root value under a special key to work around this fact.

TODO Actually use this illustration for something, should probably work with the actual example provided.



Jolie natively supports a variety of techniques for composition of services. The most important (for this work), which we will cover here are **aggregation** and **embedding**.

Embedding allows for a larger service to run smaller services as inner components. These services communicate with each other using more efficient local communication. These embedded services can be other Jolie services, but may also be services written in, for example, Java or even JavaScript. Communicating with these services is done exactly the same way as with any other service, it is entirely transparent to the application code where the service is located. TODO Can probably say a few more words about this subject. Might be easier to add this later when we know what we actually need.

Aggregation is a generalisation of proxies and load balancers. An illustration of this concept can be seen in Figure 2.1. Aggregation is useful for creating a wide variety of proxy like architectural patterns. The aggregation feature is often used along side the courier and interface extender features. Couriers allows the developer to insert code in-between the receiving the request and forwarding it. These features for example allow you to add authentication to a service which otherwise doesn't have it. This is done entirely without having to touch the original service.



Figure 2.1: Aggregation is a generalisation of proxies and load balancers

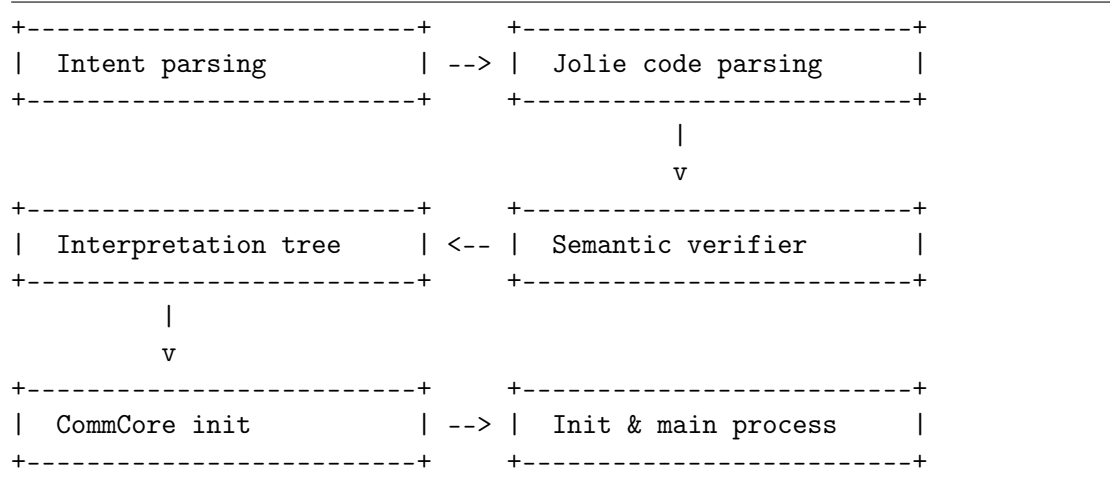
2.4 The Jolie Engine and Interpreter

- Values and types
- Phases of the compiler
- Anything we need to explain the changes that were made

In this section the internals of the Jolie engine will be introduced. This should give the reader the necessary knowledge to understand the changes made to support a module and package system for Jolie.

At the core of the Jolie engine is the interpreter. Each interpreter is responsible for parsing and executing a Jolie program. A single engine may run several instances of the interpreter, this is most commonly the case when embedding several other Jolie services.

A simplified view of the Jolie interpreter's pipeline can be seen in Figure 5.



Listing 5: A simplified view of the Jolie interpreter pipeline

The first phase of the interpreter is parsing the intent. In this phase we essentially figure out why the interpreter has been created, and what actions it should perform.

The Jolie engine can be invoked from the command-line, the command-line is the source of the intent which starts the first interpreter. The syntax for the command line is (roughly) as follows: `jolie [commands and options] <program> [program arguments]`. The options passed change the overall behaviour of the engine, and all interpreter instances share these. Commands and program arguments, however, only belong to the interpreter that they were originally passed.

The intent parsing phase is also responsible for locating and retrieving the program used. This is passed to the program parser. The program parser is the second phase, and is responsible for creating the Abstract Syntax Tree (AST) which represents the input program. The parser will produce only a single root node, namely the **Program** node. As a consequence of this, any file which is included is semantically identical to copying and pasting the source code of that file into the original file, in place of the include statement. It should also be noted that include paths are *not* relative to the file that includes, but rather relative to the current working directory (i.e. where the engine was started).

The third phase traverses the AST to make sure that it is semantically valid. This weeds out programs which are syntactically correct, but do not make sense. The amount of work done in this phase is somewhat limited, given the otherwise dynamic nature of the language.

Given a semantically valid AST the interpreter is ready to build the interpretation tree. The interpretation tree contains new nodes, known as processes. Each of these processes can be run, to execute the correct behaviour.

Once the interpretation tree is build, we're ready to execute the actual code (which lives

in the interpretation tree). First the `init` block is run, followed by the `main` block.

2.5 A Complete Application with Jolie

In this section we will describe a complete application written in Jolie. The application will contain several services, and will be written using best practices from before the module and package system.

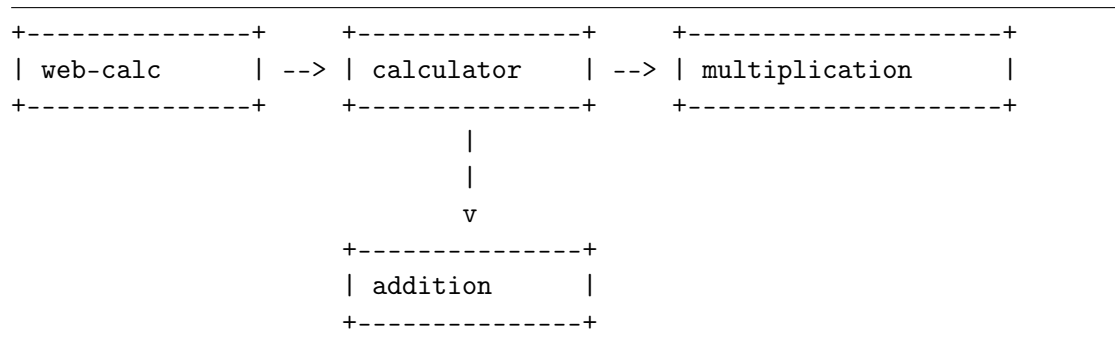
Along the way several patterns used for re-using code will be highlighted. This will serve as the core motivation for our solution.

2.5.1 Architecture

The application we will be building is a very simple calculator system. The architecture of our application is shown in Figure 6 we see an illustration of the system's architecture.

The calculator service will provide various operations for applying an operator on a sequence of numbers. In order to perform this, highly complex, action of applying these operators, the calculator will contact other microservices.

The dashed region (TODO) displays services that should be running in the same Jolie engine. This is accomplished via embedding, which was introduced in a previous section.



Listing 6: The Architecture of a Simple Microservice System

2.5.2 Implementing the Calculator Service

We will keep our focus on the *calculator* service, and its interacting with other services. Putting together the stuff learned from the previous sections, we can quickly setup an input port for the service, which has the appropriate interface. It is considered best

practice to place the public interfaces that a service exposes in its own separate file. Files intended for other services to include typically have the file extension `.iol` as opposed to `.ol`. There is no technical difference between the two, but it allows for the developer to more easily express intent. Files that are included can be placed in the `include` directory, this directory is always implicitly added to the search path.

Thus in order to implement our calculator service we create two files, one for the service implementation (`calculator.ol`), and another which can be used by other services (`calculator.iol`). The implementation and interfaces are shown in Listing 7.

<pre> 1 // calculator.ol 2 include "console.iol" 3 4 inputPort Calculator { 5 Location: "socket://localhost:12345" 6 Protocol: sodep 7 Interfaces: ICalculator 8 } 9 10 main { 11 [sum(request)(response) { 12 println@Console("Implementation 13 ↳ goes here")() 14 }] 15 [product(request)(response) { 16 println@Console("Implementation 17 ↳ goes here")() 18 }] 19 }</pre>	<pre> 1 // calculator.iol 2 type Numbers: void { 3 .numbers[2, *]: int 4 } 5 6 interface ICalculator { 7 RequestResponse: 8 sum(Numbers)(int), 9 product(Numbers)(int) 10 }</pre>
---	---

Listing 7: TODO Caption

Observation 1. The interfaces and types of a service are typically separated into their own files. This files typically has the extension `.iol` to indicate that it does not contain a service implementation, but is rather intended for inclusion by a service which requires it.

This file is typically put in the `include` directory. This directory is always added to the search path natively by the Jolie engine.

The operations that the calculator exposes, needs to collaborate with the other services. In order for us to speak to them they need output ports.

First of all the output ports needs interfaces. Like we did with the calculator service, the

other services have exposed their interfaces in a special file intended for inclusion. As a result we will have to copy the `.iol` files of these services into our own.

Secondly these output port needs to be reached. We can either embed the services, making it run inside of the same Jolie engine as our calculator service, or we can provide external bindings to it. For output ports we may change this binding dynamically at runtime. Note that this is unlike input ports which must be ready at deployment time.

Binding an output port to an external service is relatively easy. For example to let the `multiplication` port bind to a service using https, we might write `Location: "socket://mult.example.co` and `Protocol: https`. The input port at the multiplication processor would also have to match this, to ensure it runs on the correct port and speaks the correct protocol.

It is a similar to bind an output port to an embedded service. However this is done by setting the `Location` or `Protocol` attributes. We must instead instruct the engine to embed the service, which most importantly requires us to point to some executable service. The Jolie engine supports several language for these embedded services, including Jolie, Java, and JavaScript. For our desired deployment, we wanted to embed the addition service inside of the calculator service. Assuming that the addition service is written as a Jolie service, and its service is implemented in `addition.ol`, then we may create an embedding as shown in Listing 9. Just like in the case of the external services, the input port of the receiving service *must match*. In the case of embedded service there must be an input port listening on the `"local"` location.

Quite often the location of an input port is considered a deployment problem. We see this quite clearly in the case of embedding a service. All current solutions in Jolie, require us to *modify the source code of a service*, simply to change where the service should listen. Best practices in Jolie attempt to make this less of a problem by including a configuration file which contains constants. The addition service, might include a file called `addition.iol` with constants setting up the location and protocol of the service. An example of this is shown in Listing 8.

<pre> 1 // addition.ol 2 include "addition.iol" 3 4 inputPort Addition { 5 Location: ADDITION_LOCATION 6 Protocol: ADDITION_PROTOCOL 7 }</pre>	<pre> 1 // addition.iol 2 constants { 3 ADDITION_LOCATION = "local" 4 ADDITION_PROTOCOL = sodep 5 }</pre>
--	---

Listing 8: A common Jolie practice for solving configuration of a service, is to include a file containing constants with the desired configuration.

Observation 2. Most configuration of Jolie services are done via constants. Jolie constants, different from most other languages, can be simple literal types, or they may even contain identifiers. This makes them rather versatile in what they may configure.

The constants are placed in a separate source file, which is included by the service requiring the configuration.

```
1 embedded {  
2     Jolie:  
3         "addition.ol" in Addition  
4 }
```

Listing 9: Embedding the `addition` service in the `Addition` output port

With the code from Listing 7 where the output port `Console` is defined. The output port points to an embedding of the console service, and is included directly in the `console.iol` file. This is a fairly common pattern used in Jolie, especially for services that work in a library-like fashion (i.e. not intended as a stand-alone service). This pattern is used for almost every single service in the Jolie standard library.

Observation 3. Services intended to be used as libraries are often contained in a single `.iol` file. This file contain everything required to set it up, included interfaces, types and an embedded output port.

This makes it very easy to use the service, however it also makes it impossible to bind to this service externally, without first getting an embedding. This isn't possible since we cannot include the interfaces and types by themselves.

With the output ports correctly configured, we may now implement the actual business logic for our calculator. For completeness sake this might look like shown in Listing 10.

```

1 [sum(request)(response) {
2     total = 0;
3     for (i = 0, i < #request.numbers, i++) {
4         // Add current number with total, and store result in total
5         add@Addition({
6             .a = total,
7             .b = request.numbers[i]
8         })(total)
9     };
10    response = total
11 }]

```

Listing 10: Implementing the checkout operation

Finally we want to reflect slightly on the file structure that this service ended up having. In order to use external services, we had to include files which contain these interfaces. This version is worse when an embedding is desired, since the entire service with its implementation is suddenly required.

The files from these external service are also hard to manage. It isn't possible to simply move the source code of a service into its own directory. This is not possible since includes, unlike in most other languages, are not relative to the file performing the include, but rather relative to the project level root. Thus if we wish to move the service to a new directory, all includes in the source code would have to be updated. A common result of this is that every single file gets dumped into the project level root. When all files are stored in the same directory, we also get a much larger chance of having name conflicts. As a result names tend to become rather large, to make it less likely that a conflict occurs.

The final file structure of the calculator service is shown in Listing 11.

```

.
+-- include
|   +-- calculator.iol
|-- addition.iol
|-- addition.ol
|-- multiplication.iol
+-- calculator.ol

```

Listing 11: File structure of the calculator service

Observation 4. In order to collaborate with other services, it is often needed to manually copy files into the service. The files required depend on if an embedding is required, when an embedding is required we will need the entire service (and any services it may depend on itself). If we just need to interface with it, we can simply include the files required for its interface files.

Structuring the files is problematic, due to `include` statements not being relative to file performing inclusion.

Jolie Modules

3.1 Summarizing the Problem

The goal of this thesis is to promote, and facilitate better code reuse in Jolie. In Section 2.5 we showed how a typical Jolie service would be written, and summarized a number of patterns we have observed.

3.2 Modules: Design and Implementation

A module in Jolie is defined as a project root directory, and optionally an entry-point. Modules are uniquely identified by a name.

The Jolie engine needs to know about these modules. The engine is informed about these modules from the intent, passed as options, which starts the engine. This information is collected in the “intent parsing” phase, and is made available to any later phase that might need it. Since the intent is passed as an option, all interpreter instances inside the engine will know about each module. As a result any embedded service will also be aware of the same modules.

As an example, we may inform the engine about a module called “foo”, which has its source code placed in `/packages/foo` with an entry-point in `/packages/foo/main.ol` we would write: `--mod foo,/packages/foo/main.ol`. This implementation technique is very close to how, for example, you would add a JAR file to the classpath of a Java program.

The result of this decision is that quite a lot of additional options may need to be passed to the engine. However this choice was purposely chosen, it is left for another tool to make this job easier. In this case a package manager is expected to take the heavy lifting, and figure out which modules exists. This allows for more freedom in how these tools are implemented, and another implementation strategy, than the one provided in the package manager, could be created without any changes to the language infrastructure. See Section TODO about how the package manager passes this information to the engine.

To allow for better organization, a new type of include has been added to the language. These include allows the developer to include from a particular package. The syntax of

this include is shown in Listing 12. When a package include is used, the search path will be altered, such that the project root is changed to that of the package (as opposed to where the engine was started). Any file included from within this package should perform its includes relative to its own root, rather than the project root.

TODO Example

The new include syntax, along with native support for modules, allows for the code to properly organized. With this a module may be placed inside of its own directory, without any of the code having to be changed.

```
1 include "<file>" from "<module>"
```

Listing 12: Extension to the include statement, made for module imports

3.2.1 Include Algorithm for Jolie

TODO To this date I still don't understand how the include algorithm works. Ask Fabrizio about this.

3.2.2 Module Include Algorithm

TODO Module includes modify the search path. This also means the injection of the include directory, which is used as a convention. This helps maintain expected behaviour.

TODO In order to implement the root switching we use a stack.

3.3 Configuration


From observation 2 we saw that most configuration was done via the inclusion of source code. This source code would expose constants (read: literal values *and* identifiers). The included source code, however, can do anything that Jolie source normally can, and isn't limited to just the desired configuration. As a result, a service developer cannot be certain that the configurator (entity who provides configuration) doesn't start messing with other details of the program. Deploying defensive programming¹ techniques against this becomes significantly more problematic, since no guarantees about the configuration source file can really be made.

¹Defensive programming techniques are usually employed for systems that require high availability, or where safety and security is required.

Distributing re-useable packages is also problematic with this approach. Several features of package management requires that the package remains read-only. For example, updating of packages require this, without it source-code merges would be required.

We also saw that a lot of issues, that should have been a deployment issue, became a code issue.

This gives us plenty of reason to explore the need for a native configuration format. Most other systems would most likely go for a system defined in user code, as opposed to natively. An example of such framework, could be Vert.x, it is a tool-kit for building reactive applications on the JVM. Examples of such “reactive applications” are microservices. The configuration workflow is shown in Figure 13. The system will retrieve, and read external configuration files, directed by the user code, and apply the configuration as needed.



Listing 13: Simplified workflow for configuration of Vert.x applications

TODO Defined deployment time somewhere

However implementing such as a system in Jolie has its problems, most of these come from the difference between general-purpose programming languages and specialized programming languages.

In general-purpose languages, the constructs (such as a server’s socket) for the microservice architecture are created in user code. As a result they are entirely accessible from user code. This make it feasible to change their behaviour, since code can run before deployment occurs.

In Jolie the constructs are managed directly by Jolie. Doing this has multiple advantages, such as less complexity in user code, but it also means that user code is capable of doing less. Jolie user code can for example not control networking directly, but is instead forced to use the abstractions provided by Jolie (sending messages). The language puts constraints on certain constructs being fully prepared directly in the source code. This is analogous to a programming language requiring the types of a struct’s field to be present

at compile time. As a result, not all constructs can be changed at run time. Concrete examples of this includes the input ports, which needs to be ready at deployment time. Thus without native support for configuration of these, it would not be possible to change the input port.

3.3.1 Introducing the Jolie Configuration Format

In the coming section we will cover the Jolie Configuration Format (file extension: `.col`). The configuration format is build to solve the issues highlighted in the previous section.

This new configuration format should allow a developer to provide configuration for a service, which works in the setting of packaged services. This will entirely replace the old system of using a mix of constants and hard-coded values.

The configuration system will only allow configuration of constructs which has been marked as such by the module. This should also allow a developer, and potentially tools, to quickly identify which constructs allow configuration. As much as possible the code should act as documentation of itself.

Additionally the system should work with the existing checking tools (e.g. the `--check` flag). The checking tools of Jolie are, for example, used by editor plugins to show errors. These checks do not execute any code. The checking algorithm should check that configuration is valid. Additional arguments can be added at a later point, which may configure if missing configuration should be considered an error.

A configuration file is build up of *configuration units*. A configuration unit is the basic entity, which encapsulates the configuration of a single Jolie module.

A unit is uniquely identified by its profile name and module. Having multiple profiles for the same module can be useful for a variety of use-cases. A common use-case, could for example be to have separate profiles for development and production.

The units hold configuration for every possible type of configurable construct in Jolie. The ones supported are:

1. Input and output ports
 - Location
 - Protocol and protocol parameters
 - Embedding of other services (output ports only)
2. General purpose parameters (constant values accessible from Jolie programs)
3. Interface rebinding

We'll introduce the configuration format through examples. The configuration format is custom, and made to mimic the syntax of Jolie.

Listing 14 shows a very simple configuration unit. This unit sets the location and protocol for the output port `A` (line 3 and 4), the location of the input port `ModuleInput` (line 8), and two parameter assignments (line 11 and 12).

```
1 profile "hello-world" configures "my-module" {
2     outputPort A {
3         Location: "socket://a.example.com:3000"
4         Protocol: sodep { .keepAlive = true }
5     },
6
7     inputPort ModuleInput {
8         Location: "socket://localhost:80"
9     },
10
11     myParameter = 42,
12     myParameter.subProperty = "hello"
13 }
```

Listing 14: A simple configuration unit named `hello-world` configuring the module `my-module`

Embedding of output ports can be performed from within a configuration unit. This moves the embedding from being a code problem to, what it should have been, a deployment problem.

Listing 15 shows the embedding of output port `A`. Note that we need to make a reference to the module, since the profile names are placed under a namespace for each module. This way multiple services can share the same name, a situation which is likely to occur with common profile names, such as “development” and “production”.

```
1 profile "hello-world" configures "my-module" {
2     outputPort A embeds "a-module" with "a-profile"
3 }
4
5 profile "a-profile" configures "a-module" {
6     // configuration of a-module goes here.
7 }
```

Listing 15: Embeddings make reference to other configuration units

As seen in the examples, it is okay to provide partial configuration as opposed to complete configuration. Certain values may be provided by the underlying module, which uses this unit. In that case the configuration unit cannot override it. In other cases partial configuration is okay due to profiles extending other profiles.

A configuration unit may extend another unit. The two, however, must configure the same module. The inheritance tree may be of an arbitrary depth, but each unit may only extend a single unit, and they must configure the same module.

If the child and parent disagree on configuration, the child always decides. For example, if unit “B” extends “A”, and they both configure the same value, then the values found in B are the ones that are correct. Listing 16 shows an example of inheritance.

```

1  profile "a" configures "a-module" {
2      aValue = 42,
3      aValue.sub = "hello",
4
5      outputPort ExternalService {
6          Location: "socket://external.example.com:42000"
7      }
8  }
9
10 profile "b" configures "a-module" extends "a" {
11     aValue = 100
12     // aValue.sub = "hello"
13     // ExternalService.location = "socket://external.example.com:42000"
14 }
```

Listing 16: Configuration units may extend other units

The module developer is often aware of what the defaults should be. For this reason default configuration profiles may be shipped along the modules, which are implicitly imported into every configuration file. The Jolie engine will look for any `.col` file in the `conf` folder. This folder should be placed relative to the module’s root. For example, if module “a” has a file called `conf/my-defaults.col`, which contains a unit called “default”. Then the user of the package may either write a configuration unit which extends this, simply by writing `profile "something" configures "a" extends "default"`, or the default directly. There is no need for any inclusion of this file.

It should be noted that no single unit is required to provide all configuration. The system doesn’t have any “abstract”² configuration units. However it is required that the configuration file provides all the necessary configuration, as declared by the module.

²As in abstract classes, a concept often used in object oriented programming

3.3.2 Input and Output Ports

Input and output ports, both existing constructs, are defined just like before in the Jolie code.

Externally configurable ports are defined by leaving out fields that should be configurable. Listing 17 shows a configurable input port.

Only the fields which are not already listed become configurable. This is useful when building a service that needs to make assumptions about its input ports. For example if a developer is building a generic web-server, it is useful to allow the developer to change the location, but the protocol should be fixed.

One typical assumption that Jolie services make about their input ports come from aliases that are made in the protocol configuration, an example is shown in line 4 of Listing 17. The pointer statement takes two variable paths, and makes the one on the left link to the one on the right. The result of this pointer alias is that any operation running from that input port will be able to set the variable `statusCode` and the status code of the HTTP request will be set appropriately. It should also be noted that the configuration units *do not* support aliases, and there are no other ways of accessing variables hidden in the protocol configuration. As a result any service which needs to do something special with its protocol configuration, like aliasing, must do it in source code. Since the aliased variable now also takes on a new semantic meaning, it also makes sense that it must be named directly in source code, and not left to configuration, where it could take any arbitrary name.

```

1 inputPort MyWebServer {
2   Protocol: http {
3     .keepAlive = true;
4     .statusCode -> statusCode;
5   }
6   Interfaces: MyWebServerInterface
7 }
```

Listing 17: A bare-bones configurable input port for a web-server

If instead, all values could be overwritten of a port then the underlying service would not be able to make assumptions about the additional features provided by, for example, its protocol. For this reason only values not specified can be overwritten.

However simply leaving out fields, and assuming they must be configurable proves somewhat problematic for output ports. Output ports require no fields, other than the interfaces, to be defined at deployment time. It is quite common for dynamic ports to not add a default location or protocol, assuming it is changed before use. To deal with this problem, it was decided to add a new keyword (`dynamic`) for output ports which

need to change their location and protocol dynamically. Only output ports marked with this keyword are allowed to be changed at run-time. Dynamic ports are also not configurable.

3.3.3 Parameters

Parameters are read-only values which are provided to a Jolie module at deployment time. A parameter is type-checked at deployment time, to ensure that its type matches what the underlying service expects. Listing 18 shows a parameter definition and its use.

The type-checking feature functions both as a mean of documentation and ensuring that the supplied configuration is valid. Like any other type-checked feature of Jolie, it is possible to opt-out simply by setting the type to be `undefined`.

<pre> 1 // service.ol 2 parameters { 3 myParameter: string { 4 .foo: int 5 .bar: bool 6 } 7 } 8 9 main { 10 println@Console(myParameter()); // 11 ↪ "Root" 12 println@Console(myParameter.foo()); // 13 ↪ 42 14 println@Console(myParameter.bar()) // 15 ↪ false 16 }</pre>	<pre> 1 // service.col 2 profile "my-profile" configures "service" { 3 myParameter = "Root", 4 myParameter.foo = 42, 5 myParameter.bar = false 6 }</pre>
---	--

Listing 18: A parameter definition and its use

From a high-level point of view, parameters are very similar to constants. However parameters weren't implemented as an extension of constants due to their technical implementation. Namely constants aren't implemented using the "Value" system, which is used by all variables and messages in Jolie, but rather implemented at the parser level.

When the Jolie parser encounters a constant definition, it will save the token from the assignment and associate it with the identifier on the left hand side. Only identifier tokens and simple value tokens are allowed. The constant definitions are limited to only

a single token. As a result it isn't possible to define more advanced tree-like values, which parameters such as `myParameter` from Listing 18 would require.

Whenever the parser reaches a places where a constant would be allowed, it will look at the next token, check if it is an identifier, and try to replace the current token with the token defined by the constant. This produces some rather surprising results and has limitations, which isn't commonly found in other programming languages. Listing 19 illustrates how the Jolie parser processes constants.

<pre> 1 constants { 2 FOO = 42 3 BAR = ActualInterface 4 } 5 6 init { 7 println@Console(FOO)() 8 } 9 10 courier Foo { 11 [interface BAR(req)(res) { 12 /* ... */ 13 }] 14 }</pre>	<pre> 1 init { 2 println@Console(42)() 3 } 4 5 courier Foo { 6 [interface ActualInterface(req)(res) { 7 /* ... */ 8 }] 9 }</pre>
--	--

Listing 19: Constants in Jolie works by replacing tokens at the parser level. Left: The input program. Right: The program which the parser ends up seeing

While it would have been possible to extend the value system to support values it was ultimately decided against. Adding both optional type-checking of constants and expanding to values was seen as too big a departure from the original intent of constants. For backward compatibility reasons constants would also still not have been pure values, but rather either an identifier or a value.

The new parameters block is an addition to the AST of Jolie programs. This addition currently only works in collaboration with the configuration and module system. The block is simply not valid to use without. If any parameter is not defined by configuration or has the wrong type the checking scripts of Jolie will throw an error.

Even though the parameters block is added to the AST it will not be visible in the final interpretation tree. We'll learn more about this in Section 3.4.

3.3.4 Interface Rebinding

TODO Need to be more explicit about which features are new, and which are existing features.

In this section we will first more closely inspect the existing aggregation, and courier features of Jolie. This leads to the introduction of interface rebinding which is explained following that.

Jolie provides the “aggregation” feature. We first introduced aggregation in the background material for Jolie. In short the feature is a generalisation of proxies. The aggregation feature simply forwards requests from one service to another. A simple proxy to the calculator service is shown in Listing 20.

```
1 include "calculator.iol" from "calculator"
2
3 inputPort Self {
4     Location: "socket://localhost:12345"
5     Protocol: sodep
6     Aggregates: Calculator
7 }
8
9 outputPort Calculator {
10     Location: "socket://calc.example.com:12345"
11     Protocol: jsonrpc
12     Interfaces: ICalculator
13 }
```

Listing 20: A calculator proxy: This service will proxy any call to the calculator service bound in the output port `Calculator`

Aggregation can be extended by using couriers, which allows for the service to run code associated with requests that are proxied. Couriers may work either for an entire interface, or any particular operation. A courier may even decide not to forward a particular call. Listing 21 shows an extension for the calculator proxy with a courier.

```

1  courier Self { // The name of a courier matches its input port
2
3      // Couriers can match an entire interface
4      [interface ICalculator(request)(response) {
5          println@Console("Received call for calculator!")();
6          forward(request)(response)
7      }]
8
9      // Or just a particular operation
10     [sum(request)(response)] {
11         // The courier may choose to forward a request, or answer the
12         // request itself
13         if (#request.numbers > 2) { forward(request)(response) }
14         else { response = request.numbers[0] + request.numbers[1] }
15     }
16 }

```

Listing 21: A courier allows additional code to run alongside a potential forwarding

Additionally Jolie supports “interface extenders”. These can, like the name suggests, extend the types of operations (in interfaces) with additional fields. These may add additional fields for every operation, or specific operations. These can also add faults to the type signature of an operation.

Like with any other output port, the service relies entirely on the interface listed in the output port to be correct. There is no communication with the target service about the correct interface. If the interface doesn’t match, it will simply fail when attempting to invoke the operation.

A consequence of this is that it isn’t possible to create entirely generic proxies without knowing the interface.

The aggregates and courier features allow for Jolie to implement many different common proxy-like patterns. These patterns mostly deal with their target service in a generic fashion, making very few, if any, assumptions about the target service. However because the aggregations feature needs the interface it would be impossible to write a fully generic proxy service.

The interface rebinding feature fixes this problem, by allowing configuration files to redefine an interface at deployment time. This way the generic service may write ordinary code, making no assumptions about the underlying service, and only at deployment time it will know which interface the target service has.

```
1 interface ITarget
```

Listing 22: Configurable interfaces are defined by leaving the body empty

The interface is bound from the configuration file, in a similar fashion to other configurable units, as shown in Listing 23.

```
1 interface ITarget = ICalculator from "calculator"
```

Listing 23: Rebinding the `ITarget` interface to the `ICalculator` interface from the calculator module

3.3.5 Syntax of the Configuration Format (COL)

In this section the formal grammar of the configuration format is presented. The grammar is presented in a ABNF-like syntax. Complete syntax is presented in Appendix 2.

All literals are case sensitive.

A configuration file starts by a list of includes, followed by a list of regions, as shown in start rule `configuration-tree`.

```
1 configuration-tree = *include *region
2
3 include = "include" qstring
4
5 region = region-header "{" region-body "}"
6 region-header = [ "profile" qstring ] "configures" qstring
7               [ "extends" qstring ]
8 region-body = *definition
9 definition = port | interface | parameter
```

The definitions inside of a region correspond to the different configurable units. The syntax is made to closely mimic the syntax of existing Jolie code.

```
1 port = input-port | output-port
2 input-port = "inputPort" identifier "{" port-body "}"
3 output-port = embedded-output-port | external-output-port
4 external-output-port = "outputPort" identifier "{" port-body "}"
5 embedded-output-port = "outputPort" identifier "embeds" qstring
6                       "with" qstring
```

```

7  port-body = *port-property
8  port-property = location-property | protocol-property
9  location-property = "Location" ":" qstring
10 protocol-property = "Protocol" ":" identifier [ protocol-config ]
11 protocol-config = inline-tree
12
13 interface = "interface" identifier "=" identifier "from" qstring
14
15 parameter = variable-path "=" value
16 variable-path = var-id *var-node
17 var-id = ( "(" qstring ")" ) | identifier
18 var-node = "." var-id [ "[" unsigned-int "]" ]
19 value = ( primitive [ inline-tree ] ) | inline-tree
20 inline-tree = "{" *(tree-child ",") tree-child "}"
21 tree-child = "." variable-path "=" value
22 primitive = qstring | int | long | double | bool

```

The values very closely resemble what Jolie allows, primarily just removing variables from the syntax. Variable identifiers can be strings, like in Jolie, to allow for special variable names which would otherwise not be valid identifiers. This is primarily useful for building dictionary like structures with arbitrary string keys.

Additional configuration file supports C-style comments, these are allowed anywhere and are simply ignored.

```

1  comment = single-line-comment | multi-line-comment
2  single-line-comment = "//" <any text except line breaks>
3  multi-line-comment = "/*" <any text except "*/"> "*/"

```

3.4 Implementing the Configuration Format

The new interpreter pipeline is as follows, new phases marked in bold text:

1. Intent parsing
2. Jolie code parsing
3. **Configurator**
4. Semantic verifier
5. Interpretation tree builder
6. **Late verification**

7. Communication core init
8. Init & main process

The primary work occurs in the configurator service, but some of the verification work, namely the parameters, are postponed until after the interpretation tree has been build.

3.4.1 The Configuration Phase

The configuration phase builds upon the result of the intent and parsing phase. From the intent phase we receive information about modules and potentially receive the intent to use a certain configuration unit. From the parsing phase we obviously get the AST.

The configuration phase works by modifying the AST to insert the configuration gathered from the configuration unit. The processed AST is handed to the later phases, and the remainder of the pipeline works mostly like it has before.

Processing of Configuration Files

Resolving configuration files The first step of the configurator is to gather a fully processed configuration unit. This process starts by looking for the configuration file which was passed from the intent. The intent passes the configuration file using `--conf <unitName> <filePath>`.

If the file path is an absolute file path, then that file will always be used. However if the file path is relative some searching is required. The base case for resolving this is to look for the configuration file relative to the current directory. However this approach stops working once modules starts embedding other modules in a static fashion.

Consider three modules, each providing one service, **A** and **B**, **C**. Module **A** depends on **B**, and **B** depends on **C**.

```

1  /* /a/main.ol */
2  outputPort B { Interfaces: BInterface }
3
4  /* /a/config.col */
5  profile "a-deployment" configures "a" {
6      outputPort B embeds "b" with "b-deployment"
7  }
8
9  profile "b-deployment" configures "b" { /* ... */ }
10
11 /* /b/main.ol */
12 outputPort C { Interfaces: CInterface }
13
14 embedded {
15     Jolie:
16         "--conf c-deployment embedded.col c.mod"
17 }
18
19 /* /b/embedded.col */
20 profile "c-deployment" configures "c" { /* ... */ }

```

Listing 24: Relevant source code for modules A and B. Each file marked by a comment, placed in a directory corresponding to their module. For example `/a/main.ol` is the `main.ol` file which belongs to the module A

Relevant code is shown in Listing 24. Assume that we start module A with `--conf a-deployment config.col a.mod`. The problem arises because the fixed embedding of module C performed in B. Because the files are resolved relative to working directory in which the engine was started, the engine would look for the configuration file in module A. To fix this, every time an embedding is performed the engine will implicitly add a `--parent` flag. This flag informs the new interpreter (of the child) who performed the embedding. If the parent flag has been given, the configuration file will be found relative to that of the parent.

Working directory, throughout this is `/a/`. In the example given above, the following will happen:

1. Start engine (from CLI) with: `--conf a-deployment config.col a.mod` File resolved to `/a/config.col` because of relative path and no parent.
2. Configuration causes embedding of `b-deployment` with:
`--parent a --conf b-deployment /a/config.col b.mod`. File resolved to `/a/config.col` because of absolute path.

3. Source embedding in B causes embedding of `c-deployment` with:
`--parent b --conf c-deployment embedded.col c.mod`. File resolved to
`/b/embedded.col` because of relative path with parent.

Parsing of configuration files With the configuration file resolved, it is passed to the configuration parser along with a list of known modules. The parser follows the conventions established by the existing Jolie code parser. The parser is a hand-written recursive descent parser with one token lookahead.

The parser directly outputs the AST. At the top level the AST consists of **Regions** which directly map to the configuration unit abstraction. Each **Region** are namespaced under the module that they configure. Internally each namespace maps the name of the **Region** to the actual **Region**. The top-level node of the AST contains a dictionary mapping between the name of modules and their namespaces. The Java type for this mapping ends up being `Map<String, Map<String, Region>>`, this allows for easy access of **Regions** when they are needed. This does however also mean that duplicates must be caught during parsing, since they would otherwise just silently override previous entries. This would typically have been performed in a later stage to keep the parser “pure”. In this case some purity was sacrificed for some efficiency and ease of use.

The grammar of the configuration format closely follows the grammar of Jolie code. As a result the configuration format also re-uses AST nodes where possible. As we will see later, this proved helpful when outputting a configured AST. The configuration grammar only supports a subset of the features that Jolie code does.

Recall from Section 3.3.1 that each module may publish default configuration units, which can be used by user-level configuration units. For this reason the parser will make special note of references to other configuration units. Such references currently only appear in output port embeddings (e.g.

```
outputPort Foo embeds "foo" with "a-foo-unit").
```

Once the initial file has been parsed, every single configuration file placed in the “conf” directory of referenced modules are parsed and placed in the same AST. Configuration files parsed during this phase may itself have references to other modules, in that case these will be pulled in as well. A set of parsed defaults is kept to avoid parsing the same file multiple times, or potentially infinite loops.

Merging configuration files With the entire configuration tree in place, merging is performed to create a single **Region** which represents the actual configuration. The merging occurs between a **child** and its **parent** region (i.e. the one it extends, if any) and results in a **merged** region. The **parent** is always merged (via a recursive call) first. The first **child** selected is the **Region** matching the profile name and package given by the intent.

For ports, the merging algorithm works by inserting all the ports of the **child** into the **merged** region. The ports of the **parent** are then merged with the preliminary output of the **merged** region. The process of merging a port of the **parent** against the **merged** region is shown, in pseudo code, in Listing 25. In short the merging will favor the **child**, and the **parent** will only provide values if the **child** doesn't.

```

1 procedure merge(port: Port, merged: ConfigTree): Region {
2     if (!merged.contains(port)) return port
3     existing = merged.getExisting(port)
4
5     if (existing.isComplete) return existing
6     // The existing port cannot be an embedding, since those cannot
7     // be partial.
8
9     for (property in Port.properties) {
10         if (existing[property] == null) {
11             existing[property] = merged[property]
12         }
13     }
14     return existing
15 }
```

Listing 25: Pseudo code for merging a **parent**'s **port** into a **merged** configuration tree

For parameters the same idea applies. However each parameter definition node in the AST potentially only provides a partial value. For example the node `foo.bar = 42`, where the type of `foo` is `string { .bar: int }`, only provides a partial definition of `foo`. For this reason the parameter nodes are first added from the **parent** and then by the **child**. This allows the definitions of the **child** to override those of the **parent**.

Applying Configuration to the AST

The **merged** region from the previous section is then used along with the AST corresponding to the Jolie program being run. The last phase of the configurator will process the entire AST and output a new AST, which is fully configured. This process works by running a **process** function on every single top-level node. The **process** function will inspect the node, and depending on the node return a list of replacement nodes.

Throughout this process the configurator will also verify that no configuration is performed on nodes which aren't configurable. This includes dynamic ports and non-empty interface definitions.

Ports When a port node is encountered in the AST, the **merged** region is searched for an output port matching its name. If none is found, the AST node is returned as it is, no modifications are made to it. Otherwise the configuration proceeds. For input ports and non-embedding output ports the process is the same. The properties of the AST port and the port represented in the configuration region are compared one-by-one. If any property is defined by both, an error is returned pointing to the AST node and the place that the configuration took place. At the end, if no errors are returned, a new port is returned with their properties merged together.

For output ports which embed another module the process is slightly different. The configurator will first verify the existence of the configuration unit by looking it up in the configuration tree. If it is not found an error will be returned. Otherwise the configurator will output a new embedding node. The embedding node will contain a reference back to the same configuration file the current **Region** comes from. When the embedding is performed, the configurator of the embedded service will lookup the correct configuration unit from the same file. This will however cause an addition parsing of the same file, but a caching layer at the configuration parser could take care of this. The embedding process is demonstrated in Listing 26.

<pre> 1 embedded { 2 Jolie: 3 "--conf foo-profile ↳ <inputConfigFile> foo.mod" in ↳ Foo 4 }</pre>	<pre> 1 outputPort Foo embeds "foo" with ↳ "foo-profile"</pre>
---	--

Listing 26: AST nodes generated, shown as code (left), by configuration (right) for an output port containing an embedding

Interface Rebinding Just like ports, when an interface definition is found the **merged** region is searched for a matching interface definition. Assuming no errors, the configurator must now replace the dummy interface definition with the real interface definition as the configuration specifies.

Quickly recapping interface rebinding in the configuration is shown in Listing 27. Note that the only information the configurator has for locating the replacement is the name of the interface and its module.

<pre> 1 interface Dummy = Concrete from "my-module"</pre>
--

Listing 27: Rebinding the interface **Dummy** to match the interface **Concrete** from the module **my-module**

The interface will be found by parsing the module's entry point. Its AST will then be searched for an interface matching the replacement. Simply replacing the dummy with its replacement, however, isn't enough. Consider, for example, the scenario shown in Listing 28. In this scenario, the processed AST will no longer be valid due to the missing types (`FooRequest` and `FooResponse`).

```

1  type FooRequest: void {
2      .child: string
3  }
4
5  type FooResponse: int
6
7  interface Concrete {
8      RequestResponse:
9          foo(FooRequest) (FooResponse)
10 }
```

Listing 28: Simply copying the interface definition is not enough, the types must also be copied

The relevant type definitions are found by iterating through every operation of the replacement interface. If the types listed in the operations are type links those are followed. This leads to a recursive algorithm, which looks through every type looking for type links, which will cause more types to be copied over. In order to avoid infinite loops a set of already visited types is kept. If a type has been seen before then its children will not be visited.

In the end, the original interface definition will be replaced, and all necessary type definitions will also be added to the program.

Parameters Parameters being a new concept in Jolie simply outputs new nodes for every single parameter assignment. This along with the definitions gathered from ordinary Jolie code parsing are processed in interpretation tree building phase.

3.4.2 Additions to the Semantic Verifier

As we have seen several syntax changes were also made to the core languages. These modifications to the syntax also slightly changed which programs were legal. As an example interfaces are now allowed to be empty at a syntactical level, to signal that they should be replaced by configuration. However having an empty interface not being replaced should still cause an error.

This means that additional checking must not be performed to ensure that programs are semantically valid. Such a phase already exists in the compiler, the semantic verifier.

The semantic verifier works by performing an AST traversal and validate the semantics of each node.

The verifier has been updated to not allow empty interfaces, this was previously enforced at a syntax level. Since the configurator runs before this phase and updates the AST, this will only cause errors for non-configured trees.

Non-dynamic ports are ensured to be constant by using features already implemented by the semantic verifier.

The semantic verifier is the last phase of the Jolie interpreter when run with the `--check` flag. As we will discuss in the coming section, it is not practical to perform type-checking of parameters at this point.

TODO Shouldn't the semantic verifier ensure that parameters are constant?

3.4.3 Additions to the Interpretation Tree Builder and Late Verification

Interpretation tree building was expanded to support parameters. All other features required no changes to AST structure, and could simply reuse existing features.

This section will discuss the technical of how the Interpretation Tree Builder (ITB) works. It will cover issues that are quite technical, and very specific to this particular implementation of the Jolie engine.

The Interpretation Tree Builder (ITB) is the last step before a Jolie program is ready to be executed. The ITB works by traversing the AST. The ITB will collect information from the AST in the form of definitions and inform the AST about their presence. These definitions include ports, interfaces, types, and code procedures.

Within code procedures the ITB will produce processes from the AST nodes which are runnable. It is from processes the meat of code execution is implemented. For example it is from processes everything from addition to operation calls are implemented.

In this phase the ITB will be gathering complete information about types. Type building is for the most part a fairly trivial one-to-one translation from the AST nodes into the internal type representation. The most important job of the ITB with respect to type building is resolving of type references. This turns references to a type name into an actual reference to a concrete internal type. Because types aren't build until this phase is what makes type-checking during semantic verification problematic. The semantic verifier would essentially have to duplicate the work of the ITB in order to perform type-

checking earlier. This also means that type-checking cannot be performed until the ITB is finished.

The ITB will need to initialize the system with certain values. For example, the ITB will initialize the `location` and `protocol` variables of output ports. These variables are local only to a single request. At this point, however, the state of the Jolie program has not yet been initialized.

In Jolie, state is associated with the current thread of execution. Jolie uses a separate thread of execution for handling every message. As a result every single request gets its own fresh state. All state is copied from the initialization execution thread. This thread performs work requested by the ITB followed by the `init` procedure. Thus right after the ITB is done, and the initialization execution thread is started, code execution starts. This includes both user code in form of Jolie, but also loading of any embedded services.

It should be noted that any process which works with state must be run on an execution thread, since it depends on the state being present on the thread.

The ITB receives two new node types from the configurator and code parser: parameter assignments, and parameter definitions.

For parameter definitions the ITB will resolve the types. Ensuring that all links are properly resolved. This will output a “processed” parameter definition which the ITB will inform the interpreter about.

In a similar fashion, the ITB will process parameter assignments and inform the interpreter. Type-checking cannot be performed at this point, due to types not being ready before the ITB finishes. Simply outputting processes for configuration will make it impossible to perform type-checking before code execution starts, due to their dependency on state.

The processed assignments contains an ordinarily processed LHS, which contains the variable path. The RHS has been processed into an expression, which can be evaluated into a native Jolie value on which type-checking can be performed.

Because of these constraints a new “late-checking phase” has been created. This phase runs right after the ITB, but before code execution starts. In this phase it is not possible to directly touch the init state. It is however possible to both create values, and evaluate expression. This phase creates a synthetic state by creating a new value, which will act as the state. Expression are evaluated against this root and later type-checked. The result of evaluating the parameter assignments are then copied into the init state, assuming, of course, that type-checking was successful.

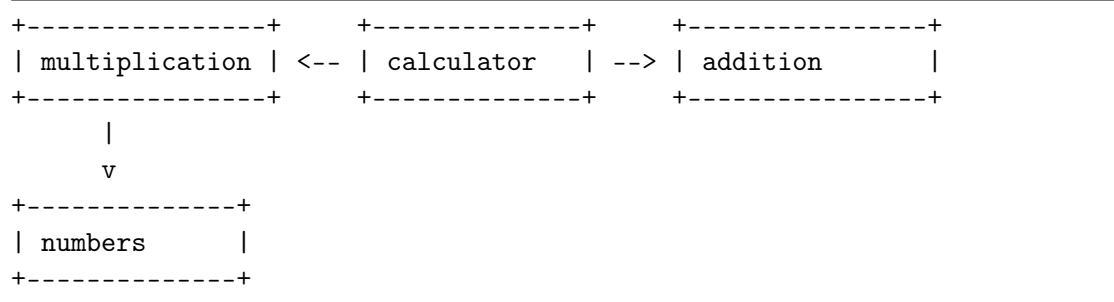
3.5 Examples

In this section we will introduce two examples, which will showcase different aspects of the Jolie module, and configuration system. These examples will be revisited again in Section TODO, where they will be shown in the context of packages.

3.5.1 Example A TODO Name

In this example we will look at the calculator example shown in the introductory chapters. The calculator system will be updated to work under the new Jolie module system. We will also show how what was previous problems that required code changes, can now be performed entirely in independent configuration files.

Quickly recapping from previous chapters. Figure 29 shows the system architecture for the calculator system. The system contains three services. The main-point of contact for a client is the `calculator` service. This service will delegate the “hard” work to the other two services, that is `addition` and `multiplication`.



Listing 29: System architecture for the “calculator system”

The file structure of our services contained a lot of duplication. The file structure is shown in Listing 30. As we can see from this there is quite a bit of duplication of files, and we cannot simply copy the directories of our dependencies, as this would place the files in the wrong places. We also resort to code files for distributing default configuration profiles. Switching between embedding a service and binding to it externally will also require significant changes.

TODO This uses the updated example from e-mail. The introduction needs to be updated to use the same example.


```

// multiplication      // calculator      // addition
.
|-- multiplication.ol  |-- addition_config.iol  .
|-- include            |-- addition.iol    |-- include
|   |-- mult_config.iol |-- addition.ol          |   |-- addition_config.iol
|   |-- multiplication.iol |-- calculator.ol        |   |-- addition.iol
|-- numbers.iol        |-- include            |-- addition.ol
                        |   |-- calculator.iol
// numbers             |-- mult_config.iol
.                       |-- multiplication.iol
|-- include            |-- multiplication.ol
|   |-- numbers.iol    |-- numbers.iol
|-- main.ol            |-- numbers.ol

```

Listing 30: Initial file structure of the calculator system

The first step is to update all includes which require files from other modules, to use the new module includes. For example, `multiplication.iol` depend on types from the utility module `numbers` from the file `numbers.iol`. Their includes change from the old `include "numbers.iol"` to the new module include `include "numbers.iol" from "numbers"`.

Using module includes also provides a certain level of namespacing of file names. This system will allow for two files to be named the same thing, without causing problems. This allows us to give files better names, and allow many rather useless prefixes. For example the file `addition.iol` can be renamed to a more meaningful name such as `types.iol`. Causing includes to say `include "interface.iol" from "addition"`.

The next step is to convert the legacy configuration into the new configuration format. The contents `mult_config.iol` previously contained constants for configuring the input port and parameters of the service. Legacy configuration is shown in Listing 31.

In order to convert port, we must first decide which ports need to be configured, and which fields should be configurable. In this case we should only configure our input port. Since the module makes no assumptions about this service's input port, we can simply allow it all to be configurable.

For parameters we must determine the actual type each parameter can take. This must be put in the `parameters` block. The complete transformation is shown in Listing 32.

```
1 constants {
2     MULT_INPUT_LOCATION = "socket://localhost:12345",
3     MULT_INPUT_PROTOCOL = sodep,
4     MULT_MAX_INPUT = 1000,
5     MULT_MIN_INPUT = 0,
6     MULT_DEBUG = false
7 }
```

Listing 31: Legacy configuration file for the multiplication module

```
1 inputPort Multiplication {
2     Interfaces: IMultiplication
3 }
4
5 parameters {
6     minInput: int,
7     maxInput: int,
8     debug: bool
9 }
```

Listing 32: Updated configuration for the multiplication module

Although we're not quite finished yet with the configuration. The original service provided default configuration. A new directory named `conf` is created. Inside which we can make as many defaults as we wish. For now we can create a file named `default.col`. In Listing 33 we show some default configuration. The inheritance feature allows us to easily create profiles for various deployment configurations which a user of the module might wish to use. The `"self-hosted"` profile would be ideal for hosting the service standalone, while the `"embedded"` and `"debug"` profiles would be ideal for development.

```

1  profile "default" configures "multiplication" {
2      minInput = 0,
3      maxInput = 1000,
4      debug = false
5  }
6
7  profile "self-hosted" configures "multiplication" extends "default" {
8      inputPort Multiplication {
9          Location: "socket://localhost:12345"
10         Protocol: sodep
11     }
12
13     outputPort Numbers {
14         Location: "socket://numbers.example.com:22000"
15         Protocol: sodep
16     }
17 }
18
19 profile "embedded" configures "multiplication" extends "default" {
20     inputPort Multiplication {
21         Location: "local"
22         Protocol: sodep
23     },
24
25     outputPort Numbers embeds "numbers" with "default"
26 }
27
28 profile "debug" configures "multiplication" extends "embedded" {
29     debug = true
30 }

```

Listing 33: Providing defaults for a Jolie module. Using inheritance we can create suitable profiles for various typical deployment configurations.

At this point we can now clean-up in the directories of each module. The Jolie module system technical allows for any organization of modules, but for now we will simply put the source code of each module inside a directory called `modules`. In this directory we will place a directory named the same as the module, containing the complete source code of the modules.

After doing this, and renaming files to give them more meaningful names, we end up with the file structure shown in listing 34.

```

// multiplication          // calculator          // addition
.                          .                      .
|-- conf                  |-- include             |-- conf
|  `-- default.col       |  `-- interface.iol    |  `-- default.col
|-- include               |-- main.ol             |-- include
|  `-- interface.iol     |-- modules             |  `-- interface.iol
|-- main.ol               |-- addition            |-- main.ol
`-- modules               |-- multiplication      `--
    `-- numbers           |-- numbers
                           `--

// numbers
.
|-- include
|  `-- interface.iol
`-- main.ol

```

Listing 34: Final file structure of the three services for the calculator system

This file structure, as we shall also discuss later, is much better for potential package managers. The overhead which is associated with actually launching the service now is bigger. Currently in order to launch this program the following would be needed:

```

dan@host:/calculator $ jolie \
  --mod addition,modules/addition,main.ol \
  --mod multiplication,modules/multiplication,main.ol \
  --mod numbers,modules/numbers,main.ol \
  --mod calculator,.,main.ol \
  calculator.mod

```

In contrast this is quite a bit more than previously:

```

dan@host:/calculator $ jolie calculator.ol

```

However, most of this disappears when using a package manager or similar tool. In fact it will be as simple as:

```

dan@host:/calculator $ jpm start

```

This is also similar to how, for example, the Java ecosystem ordinarily does its additions of libraries. From the CLI this will require manually listing every JAR dependency. However such projects are usually managed through more advanced build tools, such as: Maven, Gradle or Ant.

The configuration system also helps significantly with changing between various different deployment configuration, which switch between embedding services and having them be external.

Before the difference between embedding a service, and not embedding a service would require drastically different work. The first step here would be to make sure that all files were included in the directory. Since there were no module system, these files would have to live in the same directory as all the other code, or alternatively rewrite all includes in the dependency. Secondly it would require changing all relevant configuration, and inserting an embedded section inside of our source code.

With the module system this will require absolutely no code changes. The calculator module would simply have to define the output ports for `addition` and `multiplication` as configurable. This is shown in Listing 35.

```
1 outputPort Addition      { Interfaces: IAddition      }
2 outputPort Multiplication { Interfaces: IMultiplication }
```

Listing 35: Definitions of the output ports for `addition` and `multiplication`

Creating configurations for embedding these and for externally binding is trivial as shown in Listing 36. No changes to the existing code is required at all. These profiles can safely use the `"debug"` profiles created in the default configuration files of both `addition` and `multiplication` since those are included by default.

```
1 profile "calculator-with-embedding" configures "calculator" {
2     outputPort Addition embeds "addition" with "debug"
3     outputPort Multiplication embeds "multiplication" with "debug"
4 }
5
6 profile "calculator-with-mixed" configures "calculator" {
7     outputPort Addition embeds "addition" with "debug"
8     outputPort Multiplication {
9         Location: "socket://mult.example.com:12345"
10        Protocol: sodep
11    }
12 }
```

Listing 36: Creating vastly different deployments no longer require changes to the code. They can instead be solved entirely from configuration files

Jolie Packages

4.1 Introduction

A Jolie Package is an extension of a Jolie Module. Recall that a Jolie Module was defined as a collection of resources, a name, and optionally an entry-point for the module. A package extends this concept by adding information required for package management.

A Jolie Package is described by a package manifest. The package manifest is a JSON file, which is always placed at the root of the package, and must be called `package.json`. The fixed location allows for the package manager to easily identify a package. The JSON format was chosen as it plain-text, and easy to both read and write for both humans and machines.

In listing 37 we show a simple package manifest. This manifest showcases the most important features of the manifest. A complete specification of the package manifest format can be seen in Appendix 1. The service that this manifest describes is shown in figure 29.

```
1 {
2   "name": "calculator",
3   "main": "main.ol",
4   "description": "A simple calculator service",
5   "authors": ["Dan Sebastian Thrane <dathr12@student.sdu.dk>"],
6   "license": "MIT",
7   "version": "1.0.0",
8   "dependencies": [
9     { "name": "addition", "version": "1.2.X" },
10    { "name": "multiplication", "version": "2.1.0" },
11    { "name": "numbers", "version": "1.0.0" }
12  ]
13 }
```

Listing 37: A Simple Package Manifest

This manifest configures the ongoing example of the calculator system. Lines 2-3 take care of the module definition. The remaining attributes, however, are entirely unique to packages. Some attributes included in the manifest are there for indexing and discoverability purposes, examples of such attributes are shown in lines 4-6. The rest of the manifest describes the dependencies of this package.


4.2 Architecture

Before continuing with the packages we'll take a quick de-tour to look at the overall architecture of JPM. This should help the reader understand in which context the packages exist. The details of this architecture will be covered in Chapter 5.

The entirety of the ecosystem around JPM is written in Jolie, using the features that the Jolie Module System provides, along with the features that JPM itself provides.

At the ten-thousand foot view of the architecture, it consists of three core services, as shown in figure 38:

1. **Registry**: Responsible for serving packages known to the registry.
2. **JPM**: Provides the back-end of JPM. This includes communication with one, or more, registries, for example to download packages.
3. **CLI**: Provides the front-end of JPM. The front-end is responsible for displaying a user-facing interface, and will communicate with the back-end to perform the actual work.



Listing 38: Ten-thousand foot view of the JPM architecture

4.3 Package Manifest

In this section we'll cover, in details, the format of Jolie manifests. The complete specification for the package manifest can be found in Appendix 1. However in this section we'll also cover some of the reasoning behind the choices found in Appendix 1.

4.3.1 The name Attribute

The **name** attribute uniquely identifies a package. We'll start this discussion by looking at the constraints imposed on it, from Appendix 1:

1. The name of a package is *not* case-sensitive
2. The length of a name is less than 255 characters
3. Names may only contain unreserved URI characters (Section 2.3 of [RFC3986])
4. Names are US-ASCII
5. Every registry must only contains a single package with a given name.

Points 1 through 4 all contribute towards a common goal: packages should be displayable in URLs. This will allow us to more easily create, for example, a web application where a developer can browse through packages.

Putting some constraints on the package name is also helpful in multiple places of the package manager. Consider for example a system that might store packages organized by their name (say for example a caching mechanism). If one package would be named “../foo”, then that might accidentally override the contents of the package named “foo”. Of course these name constraints are not a general solution for these types of injection attacks, but it does allow for several services working with packages to not worry as badly about the package names.

It should also be noted that the package names are not compatible with the rules of identifiers. This has affected the design of the Jolie module system. This is, for example, the reason that strings are used in both configuration files and module includes, when making references to other modules.

Rule 4 states that package names are unique within a single registry. But this constraint actually goes even further, since these packages directly extend from the modules, we cannot have a system containing two packages named the same thing, even if they reside in two different registries. Generally we make no assumptions that registries communicate with each other, for this reason it is perfectly allowed that two distinct registries may contain name clashes, but any single package must not have a dependency which clashes either with another dependency, or the package who has the dependency itself.

4.3.2 The Meta Data Attributes (license, authors, description)

These attributes are mostly used for search results. They allow for a developer seeking information about a package, to quickly gather the most important information.

License identifiers are validated against a pre-existing list of license identifiers. This is done to ensure that typos are not made in license names. The identifiers used are from the SPDX license list, which contains a list of commonly found licenses. Currently only the license identifiers are supported, this could be extended to use SPDX License Expression Syntax, which allows for references to custom licenses.

4.3.3 The version Attribute

The `version` attribute used by packages declare which version of the source code the manifest describes. This field uses the Semantic Versioning (SemVer) specification for version numbers. It is a requirement to use SemVer versions, if a version is specified, which is not a valid SemVer version, then validation of the manifest will fail. Enforcing this standard allows for several other features, which will be covered in a moment.

SemVer version numbers consists of three fields, all integers, those being the *major*, *minor*, and *patch* version numbers. Additionally a version may have a pre-release label or additional meta data. Labels and build meta-data won't be covered, as they are not that important for our use-case. The (simplified) syntax of SemVer version number is shown in Listing 39.

```
<major>.<minor>.<patch>[-<label>][+<meta-data>]
```

Listing 39: Simplified syntax of a SemVer version number

Precedence of version numbers are determined by comparing each unit numerically in the order of major, minor, and finally patch. The first differentiating units determine the precedence. Thus, for example: $1.0.0 < 1.0.1 < 1.1.0 < 1.10.0 < 2.0.0$.

Semantic Versioning dictates rules for how version numbers should change, depending on how the public API changes. The rules can be summarized as:

1. The patch version number *must* be incremented, if only backwards compatible bug fixes are introduced.
2. The minor version number *must* be incremented, if backwards compatible functionality is introduced.
3. The major version number *must* be incremented, if incompatible changes are introduced.

Having a major version of 0, indicates that the product is not yet stable. At this point normal rules for incrementing version numbers do not apply. As a result version 1.0.0 indicates the first public API.

Semantic Versioning has become a commonly used standard across many different package managers, examples include Cargo¹ and NPM².

TODO Highlight benefits of using semantic versioning.

For dependencies the version field allows for a bit more flexibility using SemVer expressions, also known as SemVer ranges. These expressions allows the developer more freedom

¹<http://doc.crates.io/manifest.html>

²<https://docs.npmjs.com/getting-started/using-a-package.json>

in which version to use. In short the expressions allows the developer to both whitelist certain ranges for use, and blacklist certain ranges for use.

This provides both a convenience factor, i.e. the developer can express that she doesn't care which specific version is used, as long as it has feature X (released in some known version).

This feature also provides benefits for when the dependency tree is to be calculated. Consider, for example, two dependencies A and B which both depend on C with the same level of functionality, which was released in version 1.2.0. Assuming that expressions were not allowed, these two services would have to describe their dependency in an exact version. Service A might have been written at a time when 1.2.3 of service C was the newest, while B was written a bit later and is thus depending on version 1.2.6. The problem is now that the client depending on A, and B, cannot do so, because A and B have conflicting requirements, despite the fact that they both require the same level functionality. With SemVer expressions these services may more accurately describe what they actually depend on, service A might state that it depends on " $\geq 1.2.0$ ", while B is a bit more conservative and wants exactly version 1.2.6. In this case the package manager will be able to choose version 1.2.6, since this version fulfills the requirements of all dependencies.

Allowing for SemVer expressions to be used, is however not without its drawbacks. In Section 4.5 we cover one of these problems, and its solution.

4.3.4 The Lifetime Attributes (events, main)

These attributes control aspects of the package related to the execution of scripts, by the package manager, related to the package.

The **main** attribute controls the entry-point of the module. This is used both in informing the Jolie Engine where the entry-point is located. The attribute is also used for starting the package, we discuss this in Section TODO.

The **events** attribute defines a series of lifetime hooks, which are scripts that are run when specific lifetime events occur, this is described in Section 4.6.

4.3.5 The registries Attribute

The package manager may communicate with one or more registries, when performing its work. This may either be to install dependencies, or it could be to perform some of the more organizational commands, such as creating an account.

For this the package manager will need to know about the different registries. This is what the **registries** attribute describes, it is an array containing definitions for all known registries (known to the package).

A registry is defined by a unique (within the package) name, and a location (a URI following the same rules that Jolie does for its port locations). Listing 40 shows an example.

```
1 {  
2   "name": "my-registry",  
3   "location": "socket://registry.example.com:9999"  
4 }
```

Listing 40: A registry named `my-registry` being hosted at `registry.example.com` running on port 9999

Every package manifest has an implicit entry which is named “public”. This entry points to the public registry. Every command which needs to use a registry, will optionally receive a registry name. If this name is specified, it will look within the `registries` attribute. If no name is specified, the “public” entry will be used.

TODO Should we talk about why it is good to have multiple registries here?

4.4 Dependencies

4.5 Lock Files

Lock files are a feature designed to solve some of the drawbacks associated with allowing SemVer expressions in the version field of dependencies. To see how SemVer expressions pose a problem, we must first inspect the typical engineering process in which it is used. This feature is by no means original, prior work includes, for example, Cargo, and Yarn.

Continuous integration (CI) is a software development process. This process is used to avoid the so called “integration hell”, in which the time taken to integrate, or merge, the changes from multiple developers into a single track. The process works by, multiple times a day, integrating all these changes, this way failure in integration is found much earlier. We’ll quickly summarize one possible implementation of CI here. This is not the only way to do it, but summarizes the most important parts, which we’ll use to highlight potential problems with SemVer expressions.

For this process to work efficiently it frequently uses automated tests and builds, in order to perform this integration. The automated tests and builds are then performed by a dedicated server. The server will most likely pull the source code from a single repository, which represents the most up-to-date track of development.

When new features are developed, the developer will pull the most up-to-date source code, and develop the feature locally with this. This includes performing all of the builds and tests. Thus when the feature is pushed back onto the source control all of these tests should still pass. The role of the CI server is then to verify that the process is correctly followed.

How the build and testing is performed obviously depend on which technologies are used. A typical JPM based project might perform the following steps:

1. Pull source code from version control
2. Perform the build
 - (a) Download JPM dependencies (i.e. `jpm install`)
 - (b) Potentially build other necessary artifacts (such as Java libraries)
3. Run tests

It is in step 2(a) the need for lock files arises, the problem is most easily demonstrated by an example.

Consider the following scenario: a developer wishes to develop a new feature which uses package X of at least version 2.1.0. For this reason the developer adds the following line to the dependencies section: { `"name": "X", "version": ">=2.1.0"` }. Afterwards the developer performs an install of this package, at the moment this dependency to version 2.1.3. At this point the developer continues with developing the feature, without having to perform any additional installs of the package. Once the feature is developed, the developer performs all of the tests, and see that they pass.

At this point the code is pushed onto the source control system, where it is picked up by the CI server for testing. The CI server will then perform all of the steps listed above, including installing the dependencies anew. However, since the initial install by the developer had been performed a new version has been released, say version 2.3.0. This version still fulfills the expression listed in the dependency, and is hence chosen, since this is the best fitting package for this dependency. Remember the CI server has no way of knowing which package was originally installed, it only has the package manifest to work with.

Installing a different version of the package may however lead to the tests no longer working! As a result, we cannot guarantee what works on a developers machine, necessarily works on any other machines that might try to run it. As a result, some might choose to never use SemVer expressions, to avoid this pitfall.

Lock files offer a compromise between these two extremes. A lock file contains the exact versions that every dependency was resolved to. This file should be put into source control, thus when another machine attempts to perform the build the package manager will know which exact version should be used for the build.

The convenience factor of SemVer expressions is even still left in. When we perform the initial install, we might not care so which version to use, and the package manager will still be able to just pick the best version for us. Additionally the package manager can provide upgrade scripts, which uses these expressions as guidelines.

The lock files is placed in the package directory, and is called `jpm_lock.json`. Listing 41 shows the lock file, which would have been generated for the scenario above.

It should however be noted that fresh lock files are generated for every package. This means that the lock files are *not* used by a client package. As a result library developers should still be careful not specifying too wide expressions. This feature is only intended to guarantee that a package will have the same dependencies regardless of which machine it runs on. It is not intended as a replacement for specifying exact versions. If a package requires a specific version, then that should be reflected in the package manifest.

```
1 {
2   "_note": "Auto-generated file notice",
3   "locked": {
4     "X@>=2.1.0/RegistryName": {
5       "resolved": "2.1.3",
6       "checksum": "... "
7     }
8   }
9 }
```

Listing 41: A lock file showing that the dependency for package X of at least version 2.1.0 has been resolved to version 2.1.3

4.6 Lifetime Hooks

A lifetime hook is a script that “hooks” onto certain lifetime events that occur during normal use of the package manager. These scripts allows the package developer to augment the package manager, and potentially affect the work it performs.

This is a feature which is available in many other types of software, for example the popular version control system Git³ provides such a feature.

All hooks are performed on the client-side, and never on the server side (i.e. registries). Only the hooks of the current package will be run, the hooks of a dependency will not be executed. Hooks that run before a certain action, denoted by a **pre-** prefix, may stop the corresponding command from being executed. These scripts may do so by returning a non-zero exit code.

³<https://git-scm.com/book/en/v2/Customizing-Git-Git-Hooks>

These features become useful for augmenting the workflow with in-house conventions. This may for example include ensuring that, for example, tests pass before publishing a new version. Using a simple process interface allows for the user to write these scripts in whichever technology they choose.

The hooks are defined in the “events” section of the package manifest. Listing 42 shows an example hook, which runs before publishing which ensure that all tests pass.

```
1 {  
2     // ...  
3     "events": {  
4         "pre-publish": "./run_tests.sh"  
5     }  
6 }
```

Listing 42: Defining a lifetime hook, which runs the script `run_tests.sh` before publishing the package to the registry

The following are the hooks that JPM supports:

- **pre-start**: Script is run right before the **start** command is invoked on a service.
- **post-start**: Run right before the **start** command terminates. There is no guarantee that this script is run (i.e. if the service was forcefully terminated).
- **pre-install**: Runs before the **install** command is invoked.
- **post-install**: Runs after the **install** command has finished.
- **pre-publish**: Runs before the **publish** command is invoked.
- **post-publish**: Runs after the **publish** command has finished.

4.7 The .pkg Format

The `.pkg` file format, is the format that JPM uses for distribution of packages. Simply the file format contains a Jolie package zipped up into a single file.

Currently the contents of the zip file maps directly with the underlying package. The format is however open to extension, by allowing for new files to be introduced that can be used in various extensions. One such extension could for example be native code signing, which we discuss in more details in Section TODO. An obvious implementation strategy would involve adding the necessary files directly into the ZIP archive.

Extending the ZIP file format is a quite common occurrence. Especially when shipping what is essentially a collection of files. Popular examples of other file formats using this approach includes the JAR⁴ format, and Office Open XML⁵.

4.8 Integrity Checks

This is a section about integrity checks in JPM.

- Some background about checksums (probably)
- How we do it. Reason we don't go for something like code signing (in pkg manager)
- Code signing, and why we would prefer this.

⁴Package file format typically used for aggregating many Java class files together

⁵Developed by Microsoft, used in their office applications

Package Manager

5.1 Introduction

TODO This is just an introduction to the CLI. Re-write to match the package manager overall

The command line application serves as the user interface to JPM. The application is, perhaps not unsurprisingly, named `jpm`. The tool will be used in several examples.

The command line application is responsible for displaying a more user friendly interface to inner workings of JPM. The tool will perform almost not work by itself, but will instead delegate this to the back end (See Section TODO).

When first running the tool, the user will be welcomed with the following message:

```
JPM - The Jolie Package Manager
Version 1.0.0
```

```
Usage: jpm <COMMAND> <COMMAND-ARGUMENTS>
```

```
Command specific help: jpm help <COMMAND>
```

```
Available commands:
```

```
-----
init           Initializes a repository
search         Searches repositories for a package
install        Install dependencies
publish        Publish this package
start          Start this package.
```

```
[ Remaining commands removed from snippet ]
```

As clearly visible from this snippet, for the tool to do any work we must first give it something to do via a command. The commands that JPM understands almost directly mirror the functionality provided by the back end. To use JPM to create a new package, the user must simply use the `init` command. This will display a prompt, guiding the

user through the mandatory field, and automatically create a package with the required structure. This is shown in Listing 43.

```
$ jpm init
Package name
-----
> my-package

Package description
-----
> This is my package

Author: [Format: name <email> (homepage)]
-----
> Dan Sebastian Thrane <dathr12@student.sdu.dk> (github.com/DanThrane)

Private package? [Y/n]
-----
> n

$ cat my-package/package.json | json name
my-package
```

Listing 43: The jpm tool provides a user interface for common tasks. In this example, creating a new package.

5.2 Registry

Stuff we need to cover in this section:

- Core responsibilities, such as: publish and download
- Secondary responsibilities: Package information and dependencies
- Tertiary responsibilities: Users and groups
- Once responsibilities are in place we should talk about how we split up the workload.

5.3 Security

The authorization of JPM is implemented in the `security` service. The `security` service provides two different services, which work together to form the `security` service:

1. **An authentication system:** Responsible for proving identity of users.
2. **An authorization system:** Responsible for deciding which users are allowed to do what.

In this section, the client represents a service which uses the `security` services. The system is written to assume that a client *is not* an actual user. As a result most operations have been written to support various options which should not be available to actual users. This means that when the `security` service is deployed, care must be taken to ensure that it is not reachable by any ordinary user. Only the services which use the `security` service should be able to reach it. The services that use the `security` service, will then proxy only the relevant parts along.

5.3.1 Authentication

The authentication system has a user based system. Users authenticate themselves using a password. A group is a collection of users, which have an associated set of permission. A single user may be part of many groups, it is from groups that a user gains permission to use parts of the system. We cover user permissions in Section 5.3.2.

The authentication system is build following recommendations from OWASP[OWA1; OWA2; OWA3]. In this section the authentication implementation, and rules surrounding the system are summarized.

Username and Password Rules

The usernames used for the system are unique and case in-sensitive. The username is required to be between 1 and 64 characters long. All character types are allowed in usernames. Encoding of the usernames (and passwords) are deployment dependent, specifically it depends on the protocol and its configuration.

Passwords are case sensitive. The only limitation put on password is a length requirement between 5 and 128 characters.

The maximum lengths used for these are mostly to establish reasonable limits. Accepting an unlimited amount of data can be of concern for both storage and amount of network traffic. For data that should be presented, i.e. username, having a maximum is helpful for the design of user interfaces.

Storage

The usernames and passwords are stored in a SQL database. The database itself, along with its driver is configurable via the configuration system (See Section 3.3). Passwords are hashed with BCrypt¹[PM99], each individual password has its own salt. “Salting” a password is the act of adding a randomly generated string to each password. The salt itself isn’t secret, but is instead used to avoid precomputed reverse lookup tables for the hashes (rainbow tables).

Authentication Process

We start the discussion of how the authentication process works by looking at how a client will successfully authenticate itself, this is shown in Figure 5.1.

During the initial request nothing special really happens. The **security** service collaborates with both the **bcrypt** service, along with its own database. The database contains both **user** objects and **session** objects. The **user** object has already been explained, and simply contains the username, the hash of the salted password, and the salt itself. The **session** object represents an active session. It contains the actual token, which is a string generated by a cryptographically secure random number generator. Along with this token, the object contains a timestamp for generation, and a reference to the user object which created the session. The timestamp will be used for implementing timeouts for the session.

At the end of the initial request, the client receives the token. This token is what the client will have to provide for every privileged operation. An invalidation of a session may occur. This would ordinarily happen due to a logout, or potentially a timeout which is handled by the client. The **security** service also supports checking the “freshness” of a session token, by making sure it isn’t older than some amount of time. This could be useful for ensuring re-authentication before critical operations, such as changing your password.

Figure 5.1 showed us the successful case. Internally, validation will be performed at most steps. In the case of failure at any of these steps a generic error message will be returned to the client. The error message will only indicate if it was a user error or a server error. This way an attacker cannot abuse the error messages for information. For example given the error message “Invalid password” an attacker would most likely be able to assume that the username itself is correct.

¹BCrypt is a password hashing algorithm which has been used by OpenBSD and others




Figure 5.1: The authentication process when successful

5.3.2 Authorization

The authorization system based on an “Access Control Matrix”[SS94] to define its security model. An access control matrix, is a matrix A , where A_{ij} contains the *rights* that *role* i has for *resource* j .

A *right* is a single piece of information. It describes something that an entity is allowed to do for a specific resource in a role. In this implementation a right is a simple string.

A *resource* represents any entity that can have any rights associated with it. This could, for example, be a file, which might have associated rights such as “read” and “write”.

A *role* is some entity which have a set of rights for resources. In the implementation roles have been replaced with *groups*. A group is a collection of users, which are provided by the authentication system.

Table 5.1 shows an access control matrix for a few files, which can be either read from or written to.

	File 1	File 2	File 3
Group 1	read, write	read	
Group 2	read	read	read, write

Table 5.1: A sample access control matrix for some computer files, which can be read from and written to by different groups

5.3.3 Integration with JPM

The **registry** service is the only consumer of the **security** service. On top of this service the **registry** service implements the security model JPM uses. JPM supports users these map one-to-one with the user model of the **security** service. On top of this JPM supports teams which are supported by a single underlying group. The purpose of a team is to allow for multiple users to all have ownership of a single package. Every single user and team get their own group, user groups only have a single member, while a team group has members corresponding to the team. The user groups are named **users.<name>**, and teams **groups.<name>**.

When a package is downloaded from or published to the **registry** the user’s permissions will be checked. For a user to download a package, the user must be in a group having the **read** (or **write** for publishing) right for the corresponding resource. Each package has an associated resource called **packages.<name>**. A wild card resource (**packages.***) also exists. The wild card resource represents every single package, thus if a group holds the **packages.*/read** right then that group is allowed to download ever package in the **registry** . This is exactly how a **registry** might provide guest downloads.

TODO Find a place for this

Limiting number of login attempts within a period is often recommended. This is done to avoid brute-force attacks. This isn't possible in Jolie, since the sender isn't available to Jolie user code. This is most likely due to the fact that no unique sender ID exists, since Jolie can accept messages from various types of networks.

5.4 Calculating Dependency Trees

1

5.5 The JPM Back End

5.5.1 Introduction and Responsibilities

The back end service is responsible for performing most of the heavy lifting, for the commands that are sent to it from the front end. Some of the jobs are performed directly by this service, while quite a few others are delegated.

The most important partner the back end service has is the registry service. The back end will communicate with one or more registry services, which are covered more in depth in Section TODO.

The responsibilities of the back end can roughly be divided into three categories:

1. **Package management (See Section TODO)**

- Installing packages
- Publishing packages
- Upgrading packages
- Searching for packages

2. **Account management (See Section TODO)**

- Login/logout with each registry
- Create user
- Team management

3. **Helper scripts (See Section TODO)**

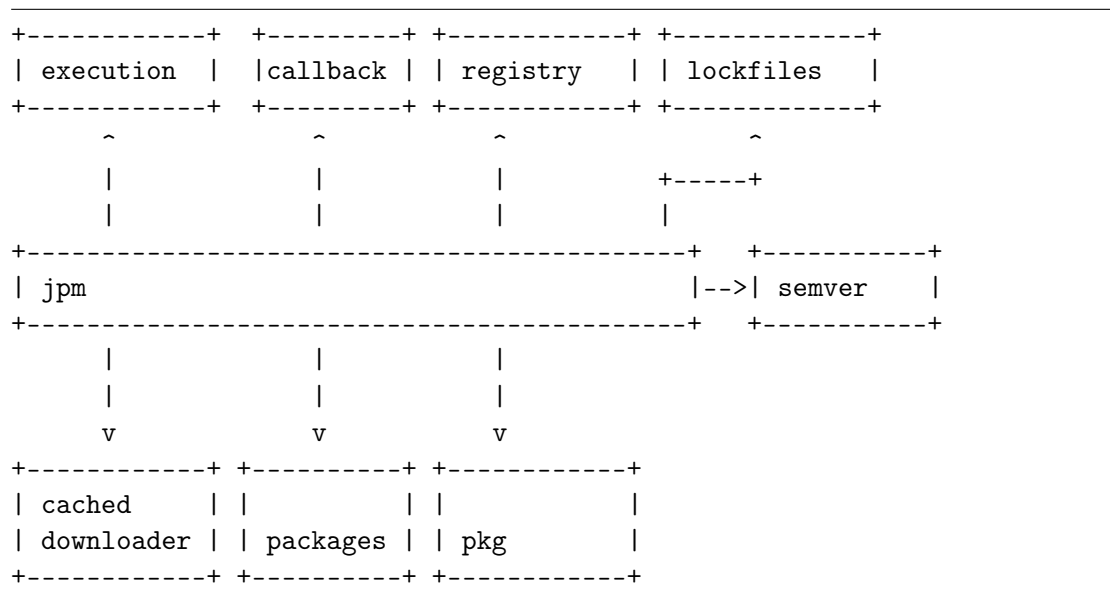
- Creating a new package

- Clearing the cache
- Starting a package

5.5.2 Architecture

The JPM back end collaborates with a few external services, these are illustrated in Figure 44.

1. **registry**: Primary collaborator, responsible for storing packages. See Section TODO
2. **execution**: Used for starting package (Section TODO), and executing life-time hooks (Section TODO)
3. **lockfiles**: Handles the lockfiles feature (Section TODO)
4. **semver**: Library package for dealing with semantic versioning (Section TODO)
5. **cached downloader**: Acts as a cache for downloaded packages (Section TODO)
6. **packages**: Validator for the internal package format (Section TODO)
7. **pkg**: Responsible for creating the binary format used for distribution (Section 4.7)



Listing 44: System architecture, from the back end's point of view

5.6 Authentication Tokens

In Section 5.3.2 we covered the authorization framework that JPM uses. When a user successfully authenticates the authorization framework will provide an authentication token proving the user's identity.

For user convenience the JPM system does not require a valid username and password combination for every privileged operation. To avoid this, JPM will store the authentication token returned by the `registry`. These are stored in a file on the file-system.

TODO I can't write smart things right now. Words no good :-)

The authentication token (or session token) is saved instead of the username and password. Since if an attacker were to gain access to both the username and password, she would be able to gain full control over the user, this includes being able to generate new authentication tokens.

If an attacker instead were to gain only the authentication token, the damage is not as bad. The attacker would still be able to pose as the user, but only while the authentication token is valid.

Several steps are taken to mitigate the damage that can be done in case of session hijacking. TODO I just called a session, without defining it.

Session timeout

Manual session invalidation (logging out)

Re-authentication before critical actions

https://www.owasp.org/index.php/Authentication_Cheat_Sheet

https://www.owasp.org/images/0/08/OWASP_SCP_Quick_Reference_Guide_v2.pdf

5.7 The JPM Cache

The `cache` of JPM is responsible for downloading packages, and maintaining the packages in a local cache.

The local cache acts as a database of known packages. This database contains both the information needed to know about the database, as well as the binary containing the package data.

The local cache consists of previously downloaded packages, along with a locally kept database which matches the type that the `registry` uses. The database kept also track the origin `registry`, since the `cache` will potentially download from multiple. The key

used for this is the location URI of each **registry** . The name cannot be used, since these differ between each package.

The primary operation that the **cache** exposes is the **installDependency** operation. This operation will install a dependency directly into a package. JPM assumes that dependencies are installed under `<pkgRoot>/jpm_packages/<depName>`. Having packages stored under a common folder makes it easy to locate. This is, for example, useful when needing to filter out the packages. Common places for this would be version control, or from JPM when publishing a package.

The **cache** will first attempt to retrieve the package from its local database, if not found it will forward the request to the registry. In order to forward the request, the operation is fed both the registry and authentication token.


When a new package is received from a **registry** , two parts will be returned: the binary itself, along with the checksum that the **registry** believes the package to have. Checksums are both checked when initially receiving the package, but also at every install from the cache. This helps prevent initial corruption and further protects from the cache itself being corrupted. A more thorough discussion of this feature can be found in Section 4.8.

The **cache** will maintain a **registry-db** just like the **registry** component does. When new entries are loaded into the **cache** , it will update the **registry-db** in a similar fashion. As a result the **cache** is technically capable of performing the exact same responsibilities as the **registry** . This could for example allow for fully offline downloads assuming the data is already loaded into the cache.

5.8 The Command Line Interface

5.8.1 Internal Organization and Deployment

The command line tool delegates most of the work to the back end service (here called **jpm**). Figure 45 shows the architecture from the CLI's point of view. Most notably is the callback server, this server is responsible for receiving information about events that occur in the back end. These are primarily used to communicate progress, this is especially useful for long running processes, such as downloading dependencies. The back end will in these cases send these events to a callback server, which can then choose to display information about this event. The need for a separate service comes mostly from a limitation in Jolie. In Jolie all communication follows either a one-way, or request-response communication pattern. As a result of this, it isn't possible for JPM to send back information while a request is being processed. To handle this the front end (in this case **jpm-cli**) will inform the back end of where it should send events.



Listing 45: The system architecture from the CLI’s point of view

Quite a lot of services exists on the client side, notable `jpm-cli` and `jpm`. From a usability perspective this is less optimal. For this reason the `jpm` binary, which ships with the package manager, will use a deployment which embeds all of the core services together. Future versions of the package manager, may wish to optionally spin up the required services, and run them as daemons². This will remove quite a bit of the overhead associated with spinning up the JVM (required for the Jolie engine) and every embedded service. Such a practice is done by similar tools, as an example the JVM based build tool Gradle provides such a daemon³. This conversion would be relatively straight forward, since the server architecture required by the daemon is already implemented.

²A daemon is a background process running on a computer

³https://docs.gradle.org/current/userguide/gradle_daemon.html

5.9 Examples and Discussion

5.9.1 The Calculator System

Continuing with the “modularized” version of the calculator system from Section 3.5.

The first step in making these usable from the package manager is of course to create package manifests. This action becomes fairly straight forward, especially since we already have a system working with the Jolie module system. The package manifest for the `multiplication` package is shown in Listing 46. The only significant multiplications here are perhaps a version number, and a formalization of dependencies (in this case the `numbers@1.0.0` package).

```

1 {
2   "name": "multiplication",
3   "description": "An multiplication package",
4   "license": "MIT",
5   "authors": ["Dan Sebastian Thrane <dathr12@student.sdu.dk>"],
6   "main": "main.ol",
7   "version": "1.0.0",
8   "dependencies": [{ "name": "numbers", "version": "1.0.0" }]
9 }
```

Listing 46: Package manifest for the `multiplication` package

The other package manifests will look fairly similar. The dependency graph ends up mimicking the system architecture exactly. However there are some subtle problems associated with the current approach.

Currently the calculator will always download the `numbers` package, despite not directly depending on it. This happens due to the dependency on `multiplication` which itself depends on `numbers`. However the `calculator` service really only depends on `numbers` if `multiplication` is embedded. At this point this may seem like a minor detail, but this can get vastly more complicated if there are more dependencies (which themselves may have dependencies and so on).

The solution to this problem, is to make sure the package manager will only download the files we actually need, and no more. This would mean if no embedding is ever performed, we will need only the dependencies required to perform interfacing. If embedding is required we will need both the interface dependencies along with a concrete implementation.

For the sake of simplicity the same dependency system is used. It is recommended that every service separately publishes its interface and concrete implementation. This approach

has the added benefit of allowing multiple implementations of same service. For closed source services a single package containing only the interface could be published.


Depending on if an embedding is desired the concrete implementation can simply be listed as a dependency, the interface dependency will implicitly be picked up (from the concrete implementation's dependencies). If an embedding is *never* desired a dependency can be made to just the interface package.

This then begs the question of when to use a concrete implementation over a dependency. TODO say something.

Additionally JPM might benefit from optional dependencies which are only when certain conditions are made, for example during some “development build”. This might prove to be a decent compromise, allowing for the convenience of embedding without directly forcing users of the package to also depend on it.

5.9.2 Using Lifetime Hooks to Improve Development Workflow

In this example we will discover how the use of Jolie lifetime hooks can improve the development workflow. In this example we will look at a hybrid Jolie-Java service. The service itself will be written in Jolie, using an internal Java service to perform some of the computational heavy-lifting. The `packages` service, which performs package manifest validation, used in the package manager is an example of such a package type. The primary flow and communication is performed entirely by Jolie, while the more “computationally heavy” parts, like parsing, are handled by the internal Java service. This architecture is illustrated in Figure 47.



Listing 47: A Jolie service using an internal Java service

The Java service uses Gradle as its build system. The build system is capable of producing a JAR file. The Jolie engine requires this file to be included in the classpath in order

to perform the embedding. If the file is placed in the `lib` directory this will happen automatically.

A small bash script was made to run JAR creation through Gradle, and then copy the artifact into the `lib` directory. Interested readers may find both the Gradle build script and bash script in Appendix 3. Both scripts are relatively simple, but should be a sufficient starting point.

```
.
|-- build
|-- java-service
|   |-- build.gradle
|   `-- src/main/java/dk/thrane/jolie/JavaService.java
|-- main.ol
`-- package.json
```

Listing 48

| Conclusion |

| Appendix |

A.1 JPM Manifest Specification

name

Field name: `name`

Optional: `false`

Type: `string`

Description: The `name` property uniquely defines a package in a registry. Every registry must only contain a single package with a given name.

Rules:

- The name of a package is *not* case-sensitive
- The length of a name is less than 255 characters
- Names are US-ASCII
- Names may only contain unreserved URI characters (see section 2.3 of RFC 3986)

If any of these rules are broken the JPM tool should complain when *any* command is invoked. Similarly a registry should reject any such package.

version

Field name: `version`

Optional: `false`

Type: `string`

Description: This property describes the current version of this package.

Rules:

- The version string must be a valid SemVer 2.0.0 string (see <http://semver.org/spec/v2.0.0.html>)

license**Field name:** `property_name`**Optional:** `false`**Type:** `string`**Description:** Describes the license that this package is under.**Rules:**

- Must be a valid identifier. See <https://spdx.org/licenses/>

authors**Field name:** `authors`**Optional:** `false`**Type:** `string|array<string>`**Description:** Describes the authors of this package**Rules:**

- The array must contain at least a single entry
- Each entry should follow this grammar:

```
name ["<" email ">"] ["(" homepage ")"]
```

private**Field name:** `private`**Optional:** `true`**Type:** `boolean`**Description:** Describes if this package should be considered private. If a package is private it cannot be published to the “public” repository.**Rules:**

- By default this property has the value of `true` to avoid accidental publishing of private packages.

main**Field name:** main**Optional:** true**Type:** string**Description:** Describes the main file of a package.**Rules:**

- The value is considered to be a relative file path from the package root.

dependencies**Field name:** dependencies**Optional:** true**Type:** array<dependency>**Description:** Contains an array of dependencies. See the “dependency” sub-section for more details.**Rules:**

- If the property is not listed, a default value of an empty array should be used

dependency**Type:** object**Description:** A dependency describes a single dependency of a package. This points to a package at a specific point on a specific registry.**dependency.name Field name:** name**Optional:** false**Type:** string**Description:** Describes the name of the dependency. This refers to the package name, as defined earlier.**Rules:** A dependency name follows the exact same rules as a package name.

dependency.version **Field name:** version

Optional: false

Type: string

Description: Describes the version to use

Rules:

- Must be a valid SemVer 2.0.0 string
- (This property follows the same rules as the package version does)

dependency.registry **Field name:** registry

Optional: true

Type: string

Description: This describes the exact registry to use. If no registry is listed the “public” registry will be used.

Rules:

- The value of this property must be a valid registry as listed in the **registries** property.

registries

Field name: registries

Optional: true

Type: array<registry>

Description: Contains an array of known registries. See the registry sub-section for more details.

Rules:

- This property contains an implicit entry which points to the public registry. This registry is named “public”.

registry

Type: object

Description: A registry describes a single JPM registry. A JPM registry is where the package manager can locate a package, and also request a specific version of a package.

registry.name Field name: name

Optional: false

Type: string

Description: This property uniquely identifies the registry.

Rules:

- A name cannot be longer than 1024 characters
- The name cannot be “public”
- No two registries may have the same name

registry.location Field name: location

Optional: false

Type: string

Description: Describes the location of the registry.

Rules:

- Must be a valid Jolie location string (e.g. “socket://localhost:8080”)

A.2 COL Grammar

Presented as an ABNF-like syntax. All literals are case sensitive.

```

1 configuration-tree = *include *region
2
3 include = "include" qstring
4
5 region = region-header "{" region-body "}"
6 region-header = [ "profile" qstring ] "configures" qstring [ "extends" qstring ]
7 region-body = *definition
8 definition = port | interface | parameter
9
10 port = input-port | output-port
11 input-port = "inputPort" identifier "{" port-body "}"

```

```

12 output-port = embedded-output-port | external-output-port
13 external-output-port = "outputPort" identifier "{" port-body "}"
14 embedded-output-port = "outputPort" identifier "embeds" qstring "with" qstring
15 port-body = *port-property
16 port-property = location-property | protocol-property
17 location-property = "Location" ":" qstring
18 protocol-property = "Protocol" ":" identifier [ protocol-config ]
19 protocol-config = inline-tree
20
21 interface = "interface" identifier "=" identifier "from" qstring
22
23 parameter = variable-path "=" value
24 variable-path = var-id *var-node
25 var-id = ( "(" qstring ")" ) | identifier
26 var-node = "." var-id [ "[" unsigned-int "]" ]
27 value = ( primitive [ inline-tree ] ) | inline-tree
28 inline-tree = "{" *(tree-child ",") tree-child "}"
29 tree-child = "." variable-path "=" value
30 primitive = qstring | int | long | double | bool
31
32 letter = <any Unicode code point recognized by Java via Character#isLetter>
33 digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
34 bool = "true" | "false"
35 int = unsigned-int | "-" unsigned-int
36 long = int ( "l" | "L" )
37 unsigned-int = 1*digit ; Jolie discards leading 0s
38 double = 1*digit "." 1*digit [ double-exp ] | 1*digit double-exp
39 double-exp = ( "e" | "E" ) [ "-" | "+" ] 1*digit
40 qstring = ( "<" *(qdtype | quoted-pair) ">" )
41 quoted-pair = "\" ( \" | \"n\" | \"t\" | \"r\" | "<" | \"u\" )
42 qdtype = <any TEXT except "<">

```

A.3 Scripts from Lifetime Hooks Example

File: build

```

1 #!/usr/bin/env bash
2 cd java-service
3 gradle jar
4 cd ..
5 mkdir -p lib

```

```
6 cp java-service/build/libs/*.jar lib
```

File: java-service/build.gradle

```
1 group 'dk.thrane.jolie'
2 version '1.0-SNAPSHOT'
3
4 apply plugin: 'java'
5
6 sourceCompatibility = 1.8
7
8 repositories {
9     mavenCentral()
10 }
11
12 dependencies {
13     testCompile group: 'junit', name: 'junit', version: '4.11'
14     compile files("/usr/lib/jolie/jolie.jar")
15 }
16
17 jar {
18     from {
19         configurations.compile.findAll { it.name != "jolie.jar" }.collect
20         ↪ { it.isDirectory() ? it : zipTree(it) }
21     }
22 }
```

Bibliography

- [MGZ14] F. Montesi, C. Guidi, and G. Zavattaro. “Service-Oriented Programming with Jolie”. In: *Web Services Foundations*. Springer, 2014, pp. 81–107.
- [MontesiW16] Fabrizio Montesi and Janine Weber. “Circuit Breakers, Discovery, and API Gateways in Microservices”. In: *CoRR* abs/1609.05830 (2016). URL: <http://arxiv.org/abs/1609.05830>.
- [Nyg07] Michael T. Nygard. *Release It!: Design and Deploy Production-Ready Software (Pragmatic Programmers)*. Pragmatic Bookshelf, 2007. ISBN: 0978739213.
- [OWA1] OWASP. *Password Storage Cheat Sheet*. https://www.owasp.org/index.php/Password_Storage_Cheat_Sheet. Accessed 22-05-17.
- [OWA2] OWASP. *Authentication Cheat Sheet*. https://www.owasp.org/index.php/Authentication_Cheat_Sheet. Accessed 22-05-17.
- [OWA3] OWASP. *OWASP Secure Coding Practices Quick Reference Guide*. https://www.owasp.org/images/0/08/OWASP_SCP_Quick_Reference_Guide_v2.pdf. Accessed 22-05-17.
- [PM99] Niels Provos and David Mazieres. “A Future-Adaptable Password Scheme.” In: *USENIX Annual Technical Conference, FREENIX Track*. 1999, pp. 81–91.
- [RFC3986] R. Fielding T. Berners-Lee and L. Masinter. *Uniform Resource Identifier (URI): Generic Syntax*. RFC 3986. Jan. 2005. URL: <https://www.ietf.org/rfc/rfc3986.txt>.
- [SS94] Ravi S Sandhu and Pierangela Samarati. “Access control: principle and practice”. In: *IEEE communications magazine* 32.9 (1994), pp. 40–48.